

Run interpreters in parallel

Language Summit 2021



Dong-hee Na



Victor Stinner



Red Hat

Use Cases

Embed Python



- vim
- Blender
- LibreOffice
- pybind11 (C++ applications)



Red Hat

Subinterpreters



- `mod_wsgi`: handle HTTP requests
- weechat plugins (IRC client written in C)



Red Hat

(A) Embed Python



- Valgrind lists leaks at exit
- `Py_Finalize()` **must release all memory** allocations done by Python.
- More important for `Py_EndInterpreter()`
- <https://bugs.python.org/issue1635741>
(created in 2007! work started in 2019)

(A) Plugin in Python



- Use subinterpreters for plugins
- IRC client written in C
- Plugin A uses Python
- Plugin B also uses Python
- Plugins are not aware of each others
- Unloading a plugin must release all memory

(B) Run in parallel



- Run multiple interpreters in parallel
- One interpreter per thread per CPU
- **N** CPUs → **N** interpreters
- multiprocessing use cases
- **Distribute** Machine Learning

(B) Single process



- A single process is more convenient and can be efficient (specific use cases)
- Admin tools are more convenient with 1 process than with N processes
- Some APIs don't work cross-processes
- Windows: creating a thread is faster than creating a process
- macOS: slow multiprocessing spawn

(B) No shared object



- Problem: concurrent access on object refcnt
- Lock or atomic operation?
→ performance bottleneck
- Pressure on few CPU cachelines for most common objects: None, True, 1, ()
- Solution: **don't share any object**, even immutable

Subinterp drawbacks



- On a crash (segfault), all interpreters are killed.
- All imported extensions must support subinterpreters.

PEPs



- PEP 384: Defining a **Stable ABI**
- PEP 489: **Multi-phase** extension module initialization
- PEP 554: **Multiple Interpreters** in the Stdlib
- PEP 573: **Module State** Access from C Extension Methods
- PEP 630: **Isolating** Extension Modules
- PEP 3121: Extension Module **Initialization** and Finalization



Red Hat

C API & extensions

Heap types



- PEP 384 & PEP 573: Support HeapType
- PyType_FromSpec()
- PyType_FromModuleAndSpec()



Red Hat

FromSpec



Modules defining heap types with PEP 384 `PyType_FromSpec()`:

- Python 3.8: 3 modules
- Python 3.9: 9 modules
- Python 3.10: 6 modules



Red Hat

FromModuleAndSpec



Modules defining heap types with PEP 573 `PyType_FromModuleAndSpec()`:

- Python 3.8: 0 modules
- Python 3.9: 2 modules
- Python 3.10: 32 modules



Red Hat

Static types



Modules still defining types as static types:

- Python 3.8: 48 modules
- Python 3.9: 39 modules
- Python 3.10: only 16 modules!

Module state



Modules using PEP 573 module state:

- Python 3.8: 8 modules
- Python 3.9: 23 modules
- Python 3.10: 48 modules!



Red Hat

Multiphase init



Modules using the PEP 489 multiphase initialization API:

- Python 3.8: 3 modules
- Python 3.9: 30 modules
- Python 3.10: 72 modules



_abc extension

_abc module example



- Convert static type into heap type
- Define module state
- Convert module to use multi-phase initialization

_abc heap type



- PyType_**Slot** and PyType_**Spec** must be implemented to define heap types.
- To manage heap type memory:
 - Py_tp_**new**
 - Py_tp_**dealloc**
 - Py_tp_**traverse** (!)
 - Py_tp_**clear**must be implemented carefully

PyType_Slot



```
static PyType_Slot
_abc_data_type_spec_slots[] = {
    {Py_tp_doc, (void *)abc_data_doc},
    {Py_tp_new, abc_data_new},
    {Py_tp_dealloc, abc_data_dealloc},
    {Py_tp_traverse, abc_data_traverse},
    {Py_tp_clear, abc_data_clear},
    {0, 0}
};
```



Red Hat

PyType_Spec



```
static PyType_Spec _abc_data_type_spec = {  
    .name = "_abc._abc_data",  
    .basicsize = sizeof(_abc_data),  
    .flags = Py_TPFLAGS_DEFAULT  
            | Py_TPFLAGS_HAVE_GC,  
    .slots = _abc_data_type_spec_slots,  
};
```



Red Hat

_abc module state



- Module state should have heap type attribute.
- Also, `get_abc_state()` API should be implemented, since accessing module state is frequently needed.

_abc module state



```
typedef struct {  
    PyTypeObject *_abc_data_type;  
    unsigned long abc_invalidation_counter;  
} _abcmodule_state;  
  
static inline _abcmodule_state*  
get_abc_state(PyObject *module)  
{  
    void *state = PyModule_GetState(module);  
    assert(state != NULL);  
    return (_abcmodule_state *)state;  
}
```



_abc PyModuleDef



- Define `_abcmodule` to implement PEP 489 **multiphase initialization**
- To manage heap type
`m_traverse`
`m_clear`
`m_free`
should be implemented carefully

_abc PyModuleDef



```
static struct PyModuleDef _abcmodule = {
    PyModuleDef_HEAD_INIT,
    .m_name = "_abc",
    .m_doc = __abc__doc__,
    .m_size = sizeof(_abcmodule_state),
    .m_methods = _abcmodule_methods,
    .m_slots = _abcmodule_slots,
    .m_traverse = _abcmodule_traverse,
    .m_clear = _abcmodule_clear,
    .m_free = _abcmodule_free,
};
```



Red Hat

_abc_exec_func



- Initialize heap type through Py_mod_exec API
- If initialization is failed return -1 if not return 0

_abc exec func



```
static int _abcmodule_exec(PyObject *module) {  
    _abcmodule_state *state;  
    state = get_abc_state(module);  
    state->abc_nvalidation_counter = 0;  
    state->_abc_data_type =  
        PyType_FromModuleAndSpec(module,  
            &_abc_data_type_spec, NULL);  
    if (state->_abc_data_type == NULL) {  
        return -1;  
    }  
    return 0;  
}
```



Get module state



Python 3.8:

```
static unsigned long abc_invalidation_counter;

static PyObject*
_abc_get_cache_token_impl(PyObject *module)
{
    return PyLong_FromUnsignedLong(
        abc_invalidation_counter);
}
```



Red Hat

Get module state



Python 3.9:

```
static PyObject*
_abc_get_cache_token_impl(PyObject *module)
{
    _abcmodule_state *state;
    state = get_abc_state(module);
    return PyLong_FromUnsignedLong(
        state->abc_invalidation_counter);
}
```



Get module state



- Python 3.8: 50.2 ns
- Python 3.10: 50.8 ns (+0.6 ns)

```
import pyperf
runner = pyperf.Runner()
runner.timeit(
    name='bench _abc',
    setup = 'import _abc',
    stmt='_abc.get_cache_token()')
```

Work done

Per-interp free lists



- float
- tuple, list, dict, slice
- frame, context, asynchronous generator
- MemoryError

Per-interp singletons



- small integer $[-5; 256]$
- empty bytes string
- empty Unicode string
- empty tuple
- single byte char (`b'\x00' – b'\xFF'`)
- single Unicode char (`U+0000 – U+00FF`)

Per-interpreter ...



- slice cache
- pending calls
- type attribute lookup cache
- interned strings:
`PyUnicode_InternInPlace()`
- identifiers: `_PyUnicode_FromId()`



Red Hat

Get empty tuple



```
static PyObject* tuple_get_empty(void)
{
    PyInterpreterState *interp;
    interp = _PyInterpreterState_GET();
    PyObject *op = interp->tuple.free_list[0];
    return Py_NewRef(op);
}
```



Per-interpreter state



- ast
- gc
- parser
- warnings

Fix daemon threads



- Random crashes at Python exit when using daemon threads
- `take_gil()` now checks in 3 places if the current thread must exit immediately (if Python exited)
- Don't read any Python internal structure after Python exited (freed memory)

Proof of Concept



- Same factorial function on **4 CPUs**
- Sequential: 1.99 sec +- 0.01 sec (ref)
- Threads: 3.15 sec +- 0.97 sec (1.5x slower)
- **Multiprocessing**: 560 ms +- 12 ms (**3.6x** faster)
- **Subinterpreters**: 583 ms +- 7 ms (**3.4x** faster)
- (Email to python-dev, May 2020)



Red Hat

TODO

TODO (easy)



- Convert remaining extensions and static types
- Make `_PyArg_Parse` per-interpreter
- GIL itself (“already done”)
- Fix unknown bugs ;-)
- Easy but lot of small issues to fix



Red Hat

TODO (hard)



- Remove static types from the public C API
- Make None, True, False singletons per interpreter
- Get the Python Thread State (tstate) from a thread local storage (TLS)

Public static types



- Remove static types from the public C API
- Replace **&PyLong_Type** with `PyLong_GetType()`
- Guido's idea: use `&PyHeapType.ht_type` for `&PyLong_Type`
- Need a **PEP** if the C API is broken.
- <https://bugs.python.org/issue40601>

None singleton



- Add an **if** to `Py_INCREF/Py_DECREF`
→ 10% slower & CPU cacheline pressure
- `#define Py_None Py_GetNone()`
→ no API issue!
- `Py_GetNone()` { return **interp**->none; }
- <https://bugs.python.org/issue39511>
Draft PR 18301

Get tstate from TLS



- `_PyThreadState_GET()` perf issue
- `C11 _Thread_local` and `<threads.h>`
`thread_local`
- x86: single MOV with FS register
- Fallback: `pthread_getspecific()`
- Function call for extensions
- <https://bugs.python.org/issue40522>
Draft PR 23976



Red Hat

Open questions



- Need a PEP for the overall isolated interpreters design.
- Extensions wrapping C libraries with shared states: need a lock somewhere.
- Another Python 2 vs Python 3 mess? No.
- Opt-in feature, adoption can be slow.

Future

Later



- API to share Python objects
- Share data and use one Python object per interpreter with locks
- Support spawning subprocesses (fork)

Questions?



Play with:

```
./configure  
  --with-experimental-isolated-subinterpreters  
#ifdef EXPERIMENTAL_ISOLATED_SUBINTERPRETERS
```

